

A Simplified Layout Algebra for GPU Kernels

Nan An
annaneuorg@gmail.com

Xiaotian Zhou
xzz5503@psu.edu

September 19, 2026

Abstract

The layout algebra behind NVIDIA’s CuTe library describes how tensor tiles map onto threads and memory. A layout is a shape with a stride for each of its entries, so the address of a coordinate is the sum of its indices times their strides. As formalized by Shah and by Cecka, its central operations, such as composition, complement, division, product, are guarded by at least six interacting admissibility conditions. We show that this complexity is an artifact of a representational choice, not of the domain.

We present a layout algebra in which every layout names the space its inputs and outputs live in: logical indices, physical indices, or thread-value indices. A layout into logical space must be a bijection from its N coordinates onto $[0, N)$. A layout into physical space must be alias-free: different coordinates land on different addresses. The exception is broadcast, where many threads read one value on purpose, and the layout has to mark where that happens. Nothing maps out of a physical index.

Composing f with g requires only that f ’s range be exactly g ’s flattened set of coordinates. To emit the index arithmetic, naively, we can always compute f , split that by division and remainder into g ’s coordinates, and feed it into g . Often the composite can be simplified to a layout as well. We show that when it can be written this way there is a coarsest such form, whose digits every other one refines, and we emit that.

A warp’s or a thread’s slice is a *restriction*. A restricted layout is no longer composable, which forces the user to settle the whole view of a tile before taking slices.

We validate the algebra by showing that CuTe’s operations, such as composition, complement, and division, can be reproduced in it. Furthermore, we provide an implementation in OCaml, and use it to rederive the atoms and usage patterns of CuTe.

The decision procedure is machine-checked. Its two ingredients — the reduction to one coordinate at a time, and the linear-time recognition of a layout on that coordinate — are proved in Coq, each as an *if and only if*, with no axioms and no admitted goals; composing them gives the procedure as a whole, so that the checks the implementation runs are proved to succeed exactly when the map is a layout over a refinement of its domain. The proved scan is extracted to OCaml and agrees with the implementation, verdict and digits alike, on every map of the enumeration we test. We also measure the system we are replacing: CuTe does not check its admissibility conditions in general, and where they fail its complement can silently violate the postconditions CuTe’s own tests assert.

The implementation is roughly 1,000 lines of OCaml and the proofs roughly 2,000 of Coq.

1 Introduction

For GPU kernels, it is important to have a scheme for describing how to move tiles of tensors between global memory, shared memory and registers. The de facto vocabulary for such schemes

is the CuTe layout algebra underlying NVIDIA's CUTLASS library [1, 2]. The most fundamental concept is the layout, which we shall reuse heavily.

Definition 1 (Layout). A *shape* is a positive integer n , or a tuple $(S_1, \dots, S_k)$ of shapes; its *size* is $\text{size}n = n$ and $\text{size}(S_1, \dots, S_k) = \prod_i \text{size}S_i$. A *coordinate* of S is an integer $i \in [0, n)$, respectively a tuple $(c_1, \dots, c_k)$ of coordinates of the S_i ; write $\text{Coord}(S)$ for the set of them. A *layout* $S:D$ is a shape S together with a *stride* D of the same structure: an integer where S is a positive integer, and a tuple $(D_1, \dots, D_k)$ where S is $(S_1, \dots, S_k)$. Its index function $\text{Coord}(S) \rightarrow \mathbb{N}$ is defined by recursion on the shape,

$$(n:d)(i) = id, \quad ((S_1, \dots, S_k):(D_1, \dots, D_k))(c_1, \dots, c_k) = \sum_{j=1}^k (S_j:D_j)(c_j).$$

Flattening is a bijection $\langle \cdot \rangle_S : \text{Coord}(S) \rightarrow [0, \text{size}S)$, taken row-major, so that the last component varies fastest, with inverse unflatten_S . A coordinate may therefore be given as a single integer, and we say that integer *is* the coordinate. The *cosize* of the layout is the least N with its image inside $[0, N)$, which is one plus the largest value it takes. The layout is *dense* when its image is an initial segment $[0, N)$ of the integers, and a *dense bijection* when it is moreover injective, in which case $N = \text{size}S$ and every value below N is taken exactly once. A *mode* is one entry of the shape together with the matching entry of the stride. Two layouts *concatenate*: for $A = S:D$ and $B = S':D'$, the layout (A, B) is $(S, S'):(D, D')$, and it sends (c, c') to $A(c) + B(c')$.

Layouts are fundamentally the correct concept here, because physical memory is linearly addressed while the logical space is multidimensional and dense as we imagine it. They are a natural way to abstract and describe those mappings.

The algebra has four operations. Three of them, composition, complement and logical division, are formalized by Shah [3]; the fourth, logical product, we take from the CUTLASS implementation. The *complement* of a layout $B = (N_0, \dots, N_\alpha):(d_0, \dots, d_\alpha)$, sorted by stride, with respect to an integer M , is

$$B_M^* = \left(d_0, \frac{d_1}{N_0 d_0}, \dots, \frac{M}{N_\alpha d_\alpha} \right) : \left(1, N_0 d_0, \dots, N_\alpha d_\alpha \right),$$

chosen so that the concatenation (B, B_M^*) has size M and its layout function is a bijection of $[0, M)$ onto itself. *Logical division* is $A / B = A \circ (B, B_{\text{size}A}^*)$, which reindexes A so that one mode runs inside the tile B selects and the other runs across copies of it. *Logical product* is $A \otimes B = (A, A_M^* \circ B)$, where $M = \text{size}A \cdot \text{cosize}B$, the smallest M for which the composition is defined. Its first mode is one copy of A , and its second walks B through the addresses A does not occupy, so the result is $\text{size}B$ copies of A placed in the pattern B names.

These operations are critical for constructing the layouts one usually wants. There is, however, a lot of complexity underneath them. Layouts are in general non-injective and non-surjective onto $[0, M)$, and CuTe always wants to write the composite as a layout again. This causes at least six condition checks. Two are in the complement, whose shape entries are quotients and so must divide: $N_{i-1}d_{i-1} \mid d_i$ for every i , and $N_\alpha d_\alpha \mid M$. The other four belong to composition. Write the flattened shape of A as $M_0 \cdots M_\alpha$. It is *left divisible* by r when $r = M_0 \cdots M_{i-1} \cdot c$ for some i , where c is smaller than M_i and divides it, and *weakly left divisible* when the same holds except that c need not divide M_i . For each mode $(N):(r)$ of B this is three checks, since the prefix must divide r , the leftover c must divide M_i , and what remains after r

must be weakly left divisible by N . One more is required across modes, that the intervals on which they act be pairwise disjoint [3].

None of these conditions appears in the CuTe API, and CUTLASS does not in general check them. Composition’s two divisibility conditions are written in its source twice over: in a strong form that is commented out, and in a weak form that runs only when both operands are statically known. The runtime assertion for the dynamic case is commented out as well, and the reason is given there — it “can cause extra registers and inappropriate compile-time/run-time failure”. The complement’s conditions are not checked at all; its shape entries are integer quotients, and what is asserted is only that they do not come out zero. The consequence is silence rather than a diagnostic. Sweeping the 675 rank-2 layouts and cotargets of Section 7, all of them static, so this is the case CuTe *can* check, the complement meets its own documented postconditions for every one of the 78 that satisfy both divisibility conditions, and fails for 40 of the 597 that do not — with no compile-time error, no assertion, and no run-time signal.

We try to resolve those complexities by restricting what a layout can be.

Explicit index spaces. The input and the output of a layout each represent one of three things: a logical index into a tile, a thread-value index, or a physical offset in memory. Therefore a layout should declare which of the three its input and its output are, and we allow only the maps worth having between them. Logical and thread-value indices map into one another in either direction. Logical indices map into physical offsets. Nothing maps out of a physical offset. By design, a thread reaches memory by composing its map into logical space with the storage map.

Requirements that follow from the space. We allow those maps and not others because the three spaces do not carry the same obligations. Logical and thread-value indices have to account for their contents losslessly, so a map into either must be a dense bijection. Composition can only happen when one layout’s image and the other’s domain agree completely, which is natural because a physical offset is forbidden as a domain. A physical offset carries far fewer requirements, since gaps in memory are common and sharing one address between threads, to read the same element, is often what is wanted, so a map into physical need only be alias-free, meaning that distinct coordinates reach distinct addresses except where the layout marks a broadcast.

Composition and simplification. The composite of two layouts is not a layout in general. But we can decide if it can be written as a layout, and when it is a layout we emit one weighted sum of the coordinates. When it is not, we emit the composition literally, at the cost of a longer expression and more work at run time.

2 Our construction of layouts

We build on top of Definition 1. First we define the spaces.

Definition 2 (Spaces). Each side of a layout carries a space from $\{\text{logical, physical, thread-value}\}$, written $A \rightarrow B$. A layout may be annotated $A \rightarrow B$ for $A, B \in \{\text{logical, thread-value}\}$, or $\text{logical} \rightarrow \text{physical}$, and no other pair is constructible. Composition takes $A \rightarrow \text{logical}$ and $\text{logical} \rightarrow C$ to $A \rightarrow C$, so $\text{thread-value} \rightarrow \text{physical}$ exists only as a composite.

Definition 3 (Replication). Extend Definition 1 by allowing the stride against a positive integer n to be the mark $\perp$ in place of an integer, with

$$(n:\perp)(i) = 0 \quad \text{for every } i \in [0, n).$$

A zero stride is rejected rather than read as replication. This is because replication is legal for reads and not for writes, since simultaneous loads of one address are a broadcast while simultaneous stores are a race. It is better not to use 0, so that the two can never be mixed.

Definition 4 (Swizzled layout). A *swizzle* is a bijection $\text{physical} \rightarrow \text{physical}$ of the form

$$\sigma(x) = x \oplus (((x \gg \text{src}) \wedge (2^b - 1)) \ll \text{dst}),$$

its two fields $[\text{src}, \text{src}+b)$ and $[\text{dst}, \text{dst}+b)$ disjoint. A layout $A \rightarrow \text{physical}$ followed by a swizzle is a *swizzled layout*, again $A \rightarrow \text{physical}$.

A swizzle runs from physical to physical, so it can only ever be the last step, and being a bijection it sends distinct addresses to distinct addresses. It therefore leaves the conditions below untouched, and matters only in the arithmetic that is finally emitted.

Definition 5 (Composition). Let $f : A \rightarrow \text{logical}$ and let $g : \text{logical} \rightarrow C$ have domain shape S_g . The *composite* $g \circ f : A \rightarrow C$ is defined when f is a dense bijection onto $[0, \text{size}S_g)$, and is the map

$$g \circ f = c \mapsto g(\text{unflatten}_{S_g}(f(c))),$$

whose domain shape is that of f .

A chain is therefore one thread-value $\rightarrow$ logical map, then any number of logical $\rightarrow$ logical maps, then one logical $\rightarrow$ physical map.

Definition 6 (Declared replicas). For a layout L with domain shape S , define a relation $\sim_L$ on $\text{Coord}(S)$:

- if L is a layout, $c \sim_L c'$ when c and c' agree at every mode whose stride is an integer, that is, they differ only at modes with stride $\perp$;
- if $L = g \circ f$ with g of domain shape S_g , then $c \sim_L c'$ when $\text{unflatten}_{S_g}(f(c)) \sim_g \text{unflatten}_{S_g}(f(c'))$.

Definition 7 (Validity). Let $L : A \rightarrow B$ be a layout with domain shape S . Then L is *valid* when

$$\begin{cases} L \text{ is a dense bijection} & \text{if } B \in \{\text{logical}, \text{thread-value}\}, \\ L(c) = L(c') \implies c \sim_L c' & \text{if } B = \text{physical}. \end{cases}$$

It is *write-valid* when $B = \text{physical}$ and in addition

$$L(c) = L(c') \implies c = c'.$$

Only a write-valid layout may be the target of a store.

Whether a layout is a dense bijection can be decided without enumeration.

Lemma 1 (Dense bijections). *Let L be a layout with no stride $\perp$, its modes carrying sizes $n_1, \dots, n_k$ and strides $s_1, \dots, s_k$ with every $n_i \geq 2$, and put $N = \prod_i n_i$. Then L is a dense bijection onto $[0, N)$ if and only if, after sorting the modes by stride, $s_{(1)} = 1$ and $s_{(j+1)} = s_{(j)} n_{(j)}$ for every j : the sorted strides are the running products of the sorted sizes.*

Proof. ($\Leftarrow$) With running-product strides the map is the flattening of Definition 1 for the shape $(n_{(k)}, \dots, n_{(1)})$, a bijection onto $[0, N)$. Positions of size 1 contribute nothing and are dropped beforehand.

($\Rightarrow$) Sort the modes by stride and write $N_j = \prod_{i \leq j} n_{(i)}$ (so $N_0 = 1$). *Claim:* $s_{(j)} \leq N_{j-1}$ for every j . Every image value below $s_{(j)}$ comes from a coordinate supported on modes $1, \dots, j-1$, since a nonzero index at any later mode contributes at least $s_{(j)}$ and no contribution is negative; those modes produce at most N_{j-1} distinct values, so surjectivity onto $[0, s_{(j)})$ forces $s_{(j)} \leq N_{j-1}$. Consequently

$$\sum_j (n_{(j)} - 1) s_{(j)} \leq \sum_j (n_{(j)} - 1) N_{j-1} = N - 1,$$

the last equality by telescoping. But the left side is the maximum of the image, and surjectivity onto $[0, N)$ requires the maximum to be exactly $N - 1$. The two sums therefore agree, and since each term on the left is at most the corresponding term on the right, they agree termwise; as $n_{(j)} - 1 \geq 1$ this gives $s_{(j)} = N_{j-1}$ for every j , which is the running-product condition. $\square$

3 Operations

We define the following operations. Each takes a layout L , and where a second argument is needed it is written out; S ranges over shapes and A over layouts.

$\text{canonical}(S)$ for a flat shape $S = (n_1, \dots, n_k)$, the layout

$$(n_1, \dots, n_k) : \left(\prod_{i>1} n_i, \prod_{i>2} n_i, \dots, n_k, 1 \right).$$

For an entry that is itself a shape, we take the canonical layout of that shape, with every stride multiplied by the product of the sizes to its right.

$$\text{canonical}((3, 4)) = (3, 4):(4, 1)$$

	0	1	2	3
0	0	1	2	3
1	4	5	6	7
2	8	9	10	11

$\text{split}(n:d, S)$ for $\text{size } S = n$, the shape S carrying the strides of $\text{canonical}(S)$ each multiplied by d , which for a flat $S = (m_1, \dots, m_r)$ is

$$(m_1, \dots, m_r) : \left(d \prod_{i>1} m_i, d \prod_{i>2} m_i, \dots, d m_r, d \right)$$

and for a nested S follows its nesting. The argument $n:\perp$ gives S with every stride $\perp$.

$$6:2$$

	0	1	2	3	4	5
0	0	2	4	6	8	10

$$\text{split}(6:2, (2, 3)) = (2, 3):(6, 2)$$

	0	1	2
0	0	2	4
1	6	8	10

$\text{divide}(L, S)$ partitions L into tiles of shape S . Let $L = (n_1, \dots, n_m):(d_1, \dots, d_m)$ and let $S = (k_1, \dots, k_m)$ with $k_i \mid n_i$. Each mode splits as

$$\text{split}(n_i:d_i, (n_i/k_i, k_i)) = (n_i/k_i, k_i):(k_i d_i, d_i),$$

and collecting the second component of every split into one group and the first into another gives

$$\text{divide}(L, S) = ((k_1, \dots, k_m), (n_1/k_1, \dots, n_m/k_m)) : ((d_1, \dots, d_m), (k_1 d_1, \dots, k_m d_m)).$$

The first group indexes a mode inside a tile and the second indexes which tile. It is easy to see that precomposing with the flattening that sends each pair (b_i, t_i) back to $c_i = b_i k_i + t_i$ recovers L .

$L = (4, 6):(6, 1)$, coloured by tile

	0	1	2	3	4	5
0	0	1	2	3	4	5
1	6	7	8	9	10	11
2	12	13	14	15	16	17
3	18	19	20	21	22	23

$\text{divide}(L, (2, 3))$: tile coordinate (t_1, t_2) as row, tile (b_1, b_2) as column

	$(0,0)$	$(0,1)$	$(1,0)$	$(1,1)$
$(0,0)$	0	3	12	15
$(0,1)$	1	4	13	16
$(0,2)$	2	5	14	17
$(1,0)$	6	9	18	21
$(1,1)$	7	10	19	22
$(1,2)$	8	11	20	23

$\text{repeat}(L, A)$ copies of L arranged in the layout of A , which is required to be a valid logical $\rightarrow$ logical layout and so a dense bijection:

$$\text{repeat}(L, A) = (S_A, S_L) : ((\text{cosize}L) D_A, D_L),$$

where $L = S_L:D_L$, $A = S_A:D_A$, and cD multiplies every stride of D by c . If L is a dense bijection then $\text{cosize}L = \text{size}L$, and the addresses $(\text{size}L)A(a) + L(v)$ take every value in $[0, \text{size}A \cdot \text{size}L)$ exactly once: the result is a dense bijection.

When L carries a swizzle σ , repeat is defined only if $\text{cosize}L$ is a power of two and both fields of σ lie below it. Then $\sigma((\text{cosize}L)a + x) = (\text{cosize}L)a + \sigma(x)$ and every copy is swizzled alike. Otherwise the copy index reaches the bits σ reads, and the operation rejects.

$\text{interleave}(L, A)$ the same copies interleaved element by element rather than laid out one after another, with A again a valid logical $\rightarrow$ logical layout:

$$\text{interleave}(L, A) = (S_A, S_L) : (D_A, (\text{size}A) D_L).$$

The coordinate (a, v) lands at $A(a) + (\text{size}A) L(v)$, so each element of L becomes a run of $\text{size}A$ consecutive addresses holding that element of every copy, where repeat instead gives each copy a run of its own. If L is a dense bijection the result is again a dense bijection onto $[0, \text{size}A \cdot \text{size}L)$. On a swizzled L it has no defined meaning and rejects.

$$L = (2, 2):(2, 1)$$

	0	1
0	0	1
1	2	3

$$\text{repeat}(L, 3:1) = (3, (2, 2)):(4, (2, 1))$$

	0	1	0	1
0	0	1	2	3
1	4	5	6	7
2	8	9	10	11

$$\text{interleave}(L, 3:1) = (3, (2, 2)):(1, (6, 3))$$

	0	1	0	1
0	0	3	6	9
1	1	4	7	10
2	2	5	8	11

$\text{broadcast}(L, S)$ $\text{size}S$ copies of L at the same addresses:

$$\text{broadcast}(L, S) = (S, S_L) : (\perp, D_L),$$

writing $\perp$ for a stride of $\perp$ at every entry of S . The coordinate (a, v) lands at $L(v)$ whatever a is. The second argument is a shape and not a layout because the copies share their addresses and so have no strides to carry. The result is valid but not write-valid.

$$\text{broadcast}(4:1, (3)) = (3, 4):(\perp, 1): \text{copy as row}$$

	0	1	2	3
0	0	1	2	3
1	0	1	2	3
2	0	1	2	3

$\text{inverse}(L)$ for L a dense bijection with modes $(n_1, \dots, n_m):(d_1, \dots, d_m)$. Give mode i the canonical weight $w_i = \prod_{j>i} n_j$ and list the modes in decreasing order of stride. Then

$$\text{inverse}(L) = (n_{(1)}, \dots, n_{(m)}) : (w_{(1)}, \dots, w_{(m)}),$$

again a dense bijection. Read both as maps on $[0, N)$, each unflattening the integer it is given into its own shape, and they are mutually inverse.

$$L = (2, 3):(1, 2), \text{ coloured by element}$$

	0	1	2
0	0	2	4
1	1	3	5

$$\text{inverse}(L) = (3, 2):(1, 3), \text{ at address } 2a + b \\ \text{the element } a + 3b$$

	0	1
0	0	3
1	1	4
2	2	5

4 Deciding whether a composite is a layout

While a composite can be natively emitted as is, sometimes we can simplify the codegen if that composite is actually itself a layout.

Fix a map $f : \text{Coord}(S) \rightarrow \mathbb{Z}$ on a shape $S = (n_1, \dots, n_m)$. We take S flat, since a nested shape has the coordinates of the flat shape of its leaves, regrouped. Write $c = (c_1, \dots, c_m)$ for a coordinate and $\mathbf{0}$ for the origin. Every layout sends $\mathbf{0}$ to 0, and composition and swizzle preserve this, so $f(\mathbf{0}) = 0$ for every composite. Maps take values in $\mathbb{Z}$ throughout this section: strides are integers, and Section 5 subtracts a constant.

Definition 8 (Refinement). A *refinement* of a flat shape $S = (n_1, \dots, n_m)$ is a shape $S' = (S_1, \dots, S_m)$ with each S_i flat and $\text{size} S_i = n_i$. A coordinate of S_i is identified with its flattening in $[0, n_i)$ (Definition 1), so $\text{Coord}(S')$ and $\text{Coord}(S)$ are the same set. A map $f : \text{Coord}(S) \rightarrow \mathbb{Z}$ is the layout $S':D'$ when

$$f(c_1, \dots, c_m) = (S':D')(\text{unflatten}_{S_1}(c_1), \dots, \text{unflatten}_{S_m}(c_m))$$

for every coordinate; for $S' = S$ this says $f = S:D$. The question of this section is whether f is a layout $S':D'$ for some refinement S' of S and some stride D' .

Lemma 2 (Separability). *Let $S = (n_1, \dots, n_m)$ be a flat shape and $f : \text{Coord}(S) \rightarrow \mathbb{Z}$. For $1 \leq i \leq m$ let e_i be the coordinate with entry i equal to 1 and every other entry 0, and define $g_i : [0, n_i) \rightarrow \mathbb{Z}$ by $g_i(v) = f(v e_i)$. Then f is a layout $S':D'$ for some refinement S' of S if and only if*

$$f(c) = \sum_{i=1}^m g_i(c_i) \quad \text{for every } c = (c_1, \dots, c_m) \in \text{Coord}(S), \quad (1)$$

and each g_i is a layout of some flat shape of size n_i .

Proof. ($\Rightarrow$) Let $f = S':(D_1, \dots, D_m)$ with $S' = (S_1, \dots, S_m)$. By Definition 1,

$$f(c) = \sum_{i=1}^m (S_i:D_i)(\text{unflatten}_{S_i}(c_i)) \quad \text{for every } c.$$

At $c = v e_i$ every term but the i th is a layout evaluated at coordinate 0, hence 0, so $g_i(v) = (S_i:D_i)(\text{unflatten}_{S_i}(v))$: the i th term is $g_i(c_i)$, which is (1), and $g_i = S_i:D_i$. ($\Leftarrow$) If $g_i = S_i:D_i$ for each i , the same identity read backwards gives $f = (S_1, \dots, S_m):(D_1, \dots, D_m)$. $\square$

Checking (1) is one pass over $\text{Coord}(S)$. What remains is one coordinate at a time: a map $g : [0, n) \rightarrow \mathbb{Z}$ with $g(0) = 0$, and the question whether there are a flat shape T of size n and a stride D with $g = (T:D) \circ \text{unflatten}_T$. To decide it we need this composite as a formula in v .

Number the entries of T from the right, $T = (m_k, \dots, m_1)$, so that indices increase with the strides below. By Definition 1, $\text{unflatten}_T(v)$ is the coordinate $(x_k, \dots, x_1)$ of T with

$$v = \sum_{j=1}^k x_j w_j, \quad w_j = m_1 m_2 \cdots m_{j-1},$$

the w_j being the strides of $\text{canonical}(T)$; so $w_1 = 1$, $w_{j+1} = w_j m_j$, and $w_k m_k = n$. Solving for the components,

$$x_j = \lfloor v/w_j \rfloor \bmod m_j,$$

which we write $d_j(v)$. Therefore

$$(T:(s_k, \dots, s_1))(\text{unflatten}_T(v)) = \sum_{j=1}^k s_j d_j(v),$$

and g is the layout $T:(s_k, \dots, s_1)$ exactly when $g(v)$ equals this sum for every $v \in [0, n)$.

Two remarks turn the search for T into a search for the w_j . An entry with $m_j = 1$ has $d_j = 0$ and contributes nothing, so we assume T has none; then $1 = w_1 < w_2 < \dots < w_k < n$, each dividing the next and w_k dividing n , which we write $1 = w_1 \mid w_2 \mid \dots \mid w_k \mid n$ and call the *chain* of T . Conversely, every such chain is the chain of exactly one such T , given by $m_j = w_{j+1}/w_j$ with $w_{k+1} = n$. Theorem 1 finds T by finding its chain.

Each component $d_j(v) = \lfloor v/w_j \rfloor \bmod m_j$ is the remainder of $\lfloor v/w_j \rfloor$ on division by m_j . The lemma rewrites $\sum_j s_j d_j(v)$ with floors alone.

Lemma 3 (Components as differences of floors). *For positive integers $w \mid w'$ and every integer $v \geq 0$,*

$$\lfloor v/w \rfloor \bmod (w'/w) = \lfloor v/w \rfloor - (w'/w) \lfloor v/w' \rfloor.$$

Let T be a flat shape of size n with k entries, with chain $1 = w_1 \mid \dots \mid w_k \mid n$, entry sizes m_j , and components d_j as above, and let integers $s_1, \dots, s_k$ and $a_1, \dots, a_k$ be related by

$$a_j = s_j - m_{j-1} s_{j-1} \quad (1 \leq j \leq k), \quad m_0 s_0 = 0,$$

a relation that determines either family from the other. Then for a map $g : [0, n) \rightarrow \mathbb{Z}$,

$$g(v) = \sum_{j=1}^k s_j d_j(v) \text{ for every } v \in [0, n) \iff g(v) = \sum_{j=1}^k a_j \lfloor v/w_j \rfloor \text{ for every } v \in [0, n),$$

and when these hold, $s_j = g(w_j)$.

Proof. Write $q = \lfloor v/w \rfloor$ and $r = w'/w$. The remainder of q on division by r is $q - r \lfloor q/r \rfloor$, and $\lfloor q/r \rfloor = \lfloor v/(wr) \rfloor = \lfloor v/w' \rfloor$ for positive integers w, r , which is the identity. With $w = w_j$, $w' = w_{j+1}$, and $w'/w = m_j$ it reads

$$d_j(v) = \lfloor v/w_j \rfloor - m_j \lfloor v/w_{j+1} \rfloor,$$

where for $j = k$, $\lfloor v/w_{k+1} \rfloor = \lfloor v/n \rfloor = 0$ on $[0, n)$. Hence

$$\sum_{j=1}^k s_j d_j(v) = \sum_{j=1}^k s_j \lfloor v/w_j \rfloor - \sum_{j=1}^{k-1} s_j m_j \lfloor v/w_{j+1} \rfloor = \sum_{j=1}^k (s_j - m_{j-1} s_{j-1}) \lfloor v/w_j \rfloor = \sum_{j=1}^k a_j \lfloor v/w_j \rfloor,$$

which is the equivalence. The relation gives $s_j = a_j + m_{j-1} s_{j-1}$, so either family determines the other. Finally, $d_l(w_j) = \lfloor w_j/w_l \rfloor \bmod m_l$ is 0 for $l > j$ because $w_j < w_l$, is 1 for $l = j$ because $m_j \geq 2$, and is 0 for $l < j$ because w_j/w_l is a multiple of $w_{l+1}/w_l = m_l$; so $g(w_j) = \sum_l s_l d_l(w_j) = s_j$. $\square$

With floors alone, first differences read the chain off g , because $\lfloor t/w \rfloor - \lfloor (t-1)/w \rfloor$ is 1 if $w \mid t$ and 0 otherwise.

Theorem 1 (Linear recognition). *Let $n \geq 2$ (for $n = 1$, $g = 0$ is the layout $(1):0$), let $g : [0, n) \rightarrow \mathbb{Z}$ with $g(0) = 0$, and let $h(t) = g(t) - g(t-1)$ for $1 \leq t < n$. Run the following scan. Keep a copy ρ of h , an integer w_{last} initially 1, and a set R initially empty. For $t = 1, 2, \dots, n-1$ in order, if $\rho(t) \neq 0$: if $t \nmid n$ or $w_{\text{last}} \nmid t$, reject; otherwise put t in R , set $a_t = \rho(t)$ and $w_{\text{last}} = t$, and replace $\rho(t')$ by $\rho(t') - a_t$ for every multiple t' of t in $[t, n)$. If no t rejects, accept.*

Then g is a layout of some flat shape of size n if and only if the scan accepts. If it accepts, $\{1\} \cup R$ is the chain of a flat shape T of size n ; g is the layout $T:(s_k, \dots, s_1)$ with $s_j = g(w_j)$; and every flat shape of size n without entries of size 1 of which g is a layout has a chain containing $\{1\} \cup R$. We call this T the coarsest. The scan takes $O(n)$ arithmetic operations.

Proof. By Lemma 3, g is a layout of the shape with chain W if and only if there are integers a_w , $w \in W$, with $g(v) = \sum_{w \in W} a_w \lfloor v/w \rfloor$ for every $v \in [0, n)$. Since $\lfloor t/w \rfloor - \lfloor (t-1)/w \rfloor = \lfloor w \mid t \rfloor$ and $g(0) = 0 = \sum_w a_w \lfloor 0/w \rfloor$, this is equivalent to

$$h(t) = \sum_{w \in W} a_w \lfloor w \mid t \rfloor \quad \text{for } 1 \leq t < n. \quad (2)$$

Suppose (2) holds for some chain W and coefficients a_w . We show by induction on t that when the scan reaches t , $R = \{w \in W : w < t, a_w \neq 0\}$ with the recorded a_w equal to those of (2), and $\rho(t) = a_t$ if $t \in W$ and $\rho(t) = 0$ otherwise. At $t = 1$, R is empty and $\rho(1) = h(1) = a_1$ because $1 \in W$ is the only divisor of 1. Given the claim at t : $\rho(t) = h(t) - \sum_{w \in R, w \mid t} a_w = \sum_{w \in W, w \mid t} a_w - \sum_{w \in W, w < t, w \mid t} a_w = \lfloor t \in W \rfloor a_t$. If $\rho(t) \neq 0$ then $t \in W$ and $a_t \neq 0$; $t \mid n$ because every element of W divides n ; and w_{last} , which is the largest element of R or else 1, lies in W below t and so divides t because W is a chain. The scan records t with a_t , and the claim holds at $t+1$; if $\rho(t) = 0$ nothing is recorded and the claim holds at $t+1$ as well. So the scan accepts, and $R \subseteq W$; with $1 \in W$, $\{1\} \cup R \subseteq W$.

Conversely suppose the scan accepts. At each t either $\rho(t)$ was 0 or $a_t = \rho(t)$ was subtracted from it, and later steps change ρ only at or beyond their own t ; so at the end $\rho \equiv 0$, that is, $h(t) = \sum_{w \in R, w \mid t} a_w$ for $1 \leq t < n$. Every element of R divides n , lies below n , and is a multiple of the element recorded before it or of 1, so $\{1\} \cup R$ is a chain; let T be its shape, and put $a_1 = 0$ if $1 \notin R$. Then (2) holds for $W = \{1\} \cup R$, so g is a layout of T by Lemma 3, with strides $s_j = g(w_j)$.

Each element of R is at least twice the one before, so the subtractions touch fewer than $n + n/2 + n/4 + \dots < 2n$ entries of ρ ; the rest of the scan is one pass over t . $\square$

The methodology is similar to Appendix A.2.1 of Axe [7].

Codegen therefore emits one of two expressions. If the map is a layout $S':D'$ for a refinement $S' = (S_1, \dots, S_m)$, it emits $\sum_{i,j} s_{ij} d_{ij}(c_i)$, with d_{ij} the components of $\text{unflatten}_{S_i}(c_i)$, for the coarsest S' ; no such expression has fewer terms. Otherwise it emits the composite as Definition 5 writes it: the value of the right operand, unflattened into the left operand's domain shape, evaluated by the left operand. Section 5 extends this to restricted maps, which do not vanish at the origin.

5 Restriction

A warp or a thread uses one slice of an address map: the map with its own warp and thread coordinates fixed.

Definition 9 (Restriction). Let $f : \text{Coord}(S) \rightarrow \mathbb{N}$ with $S = (n_1, \dots, n_m)$ flat, let $F \subseteq \{1, \dots, m\}$, and let $p_i \in [0, n_i)$ for each $i \in F$. The *restriction* of f to p is the map $f|_p$ on the coordinates of the shape $(n_i)_{i \notin F}$ given by

$$f|_p(c) = f(c'), \quad c'_i = \begin{cases} p_i & i \in F, \\ c_i & i \notin F. \end{cases}$$

$f|_p$ is in general not a layout: $f|_p(\mathbf{0})$ is the contribution of the fixed entries, which need not be 0, while every layout vanishes at $\mathbf{0}$.

$f|_p$ is in general not dense, since its image omits the other slices, so it is not a valid map into logical or thread-value (Definition 7) and cannot be the right operand of a composition (Definition 5). Restriction is therefore allowed only for maps into physical, and is applied after composition, where the warp and thread coordinates are entries of the composite's domain. Restricting a restriction is restricting by the union of the fixed entries.

The address expression of a slice is Section 4 applied to $f|_p - f|_p(\mathbf{0})$, which vanishes at the origin, plus the constant $f|_p(\mathbf{0})$; when that map is not a layout, the composite's expression is emitted with the fixed coordinates replaced by their values. The difference can take negative values, which is why Section 4 is stated over $\mathbb{Z}$: for $(2, 2):(2, 1)$ under the swizzle $\{1; 1; 0\}$, $f(a, b) = (2a + b) \oplus a$, fixing $a = 1$ gives $f|_p(b) = 3 - b$ and $f|_p - f|_p(\mathbf{0}) = -b$. The emitted arithmetic then carries a negative coefficient. For $f = (4, 8, 4):(32, 1, 8)$, $F = \{1\}$ and $p_1 = 2$, the slice is $f|_p(r, v) = 64 + r + 8v$.

6 CuTe in this algebra

CuTe lists the modes of a layout with the first varying fastest; we reverse the list, recursively, and the two then agree as functions on the same integers. CuTe's $\text{composition}(A, B)$ applies B first and is $A \circ B$ in our algebra, with B the right operand of Definition 5. We take the four operations in the order of Section 1, state what each is in our algebra, and name the documented examples [1] on which the two are checked to agree.

Composition. In CuTe, $\text{composition}(A, B)$ is the layout L with $L(c) = A(B(c))$ for every $c \in \text{Coord}(S_B)$. In our algebra, $A \circ B$ is the map $c \mapsto A(\text{unflatten}_{S_A}(B(c)))$, defined whenever B is a dense bijection onto $[0, \text{size}A)$, and Theorem 1 decides whether it is a layout. For such a B , when CuTe's composition exists the two are the same function, so Theorem 1 accepts and returns the coarsest shape that realizes it. Checked on the three documented examples,

$$\begin{aligned} \text{composition}((6, 2):(8, 2), (4, 3):(3, 1)) &= ((2, 2), 3):((24, 2), 8), \\ \text{composition}(20:2, (5, 4):(4, 1)) &= (5, 4):(8, 2), \\ \text{composition}((10, 2):(16, 4), (5, 4):(1, 5)) &= (5, (2, 2)): (16, (80, 4)), \end{aligned}$$

as functions at every coordinate; in all three B is a dense bijection onto $[0, \text{size}A)$.

Complement. In CuTe, B_M^* is the layout of Section 1, defined under two divisibility conditions. In our algebra there is no complement operator. For $B = (N_0, \dots, N_\alpha):(d_0, \dots, d_\alpha)$ the two conditions make every factor of $M = d_0 \cdot N_0 \cdot \frac{d_1}{N_0 d_0} \cdot N_1 \cdots \frac{M}{N_\alpha d_\alpha}$ an integer. Let R be the

shape of these factors in our order, $R = (\frac{M}{N_\alpha d_\alpha}, N_\alpha, \frac{d_\alpha}{N_{\alpha-1} d_{\alpha-1}}, \dots, N_0, d_0)$, and let S be R with every entry other than the N_i replaced by 1. Then

$$\text{divide}(\text{canonical}(R), S) = (B, B_M^*)$$

after dropping entries of size 1: the tile is B and the rest is its complement. Whether a given pair (B, B^*) is a bijection onto $[0, M)$ is decided by Lemma 1. Checked on the six documented complements: $\text{divide}(\text{canonical}(24), (4))$ has tile 4:1 and rest 6:4, which are $(4:1)_{24}^*$ and $(6:4)_{24}^*$; $\text{divide}(\text{canonical}((12, 2)), (4, 1))$ has tile 4:2 and rest $(3, 2):(8, 1)$, which is $(4:2)_{24}^* = (2, 3):(1, 8)$; and all six pairs (B, B_{24}^*) are bijections onto $[0, 24)$.

Division. In CuTe, $A/B = A \circ (B, B_{\text{size}A}^*)$. In our algebra, for a tile given as a shape S with $k_i \mid n_i$, this is $\text{divide}(A, S)$ of Section 3, and

$$\text{divide}(A, S) = A \circ \text{divide}(\text{canonical}(S_A), S),$$

whose right operand is a regrouping of the identity on $[0, \text{size}A)$ and so a dense bijection onto it. For a tile given as a layout B it is the composite $A \circ (B, B_{\text{size}A}^*)$ itself: $(B, B_{\text{size}A}^*)$ is a bijection onto $[0, \text{size}A)$, so Definition 5 applies with no further condition. Checked, as functions, on the documented strided tiler,

$$(4, 2, 3):(2, 1, 8) / 4:2 = ((2, 2), (2, 3)):(4, 1), (2, 8)).$$

CuTe's `zipped_divide` groups the result as $((\text{tile}), (\text{rest}))$, which is the grouping `divide` produces; `logical_divide` groups it per mode, a regrouping of the same layout.

Product. In CuTe, $A \otimes B = (A, A_M^* \circ B)$ with $M = \text{size}A \cdot \text{cosize}B$. In our algebra,

$$\text{repeat}(A, B) = (S_B, S_A):((\text{cosize}A) D_B, D_A).$$

When A is a dense bijection, $A_M^* = (\text{cosize}B):(\text{size}A)$, so

$$A \otimes B = (A, S_B:(\text{size}A) D_B) = \text{repeat}(A, B).$$

When $\text{cosize}A > \text{size}A$ the two differ. Take $A = (2, 2):(4, 1)$ and $B = 6:1$. Printed as CuTe prints layouts, one cell per coordinate holding its address and coloured by the copy it belongs to, A is

	0	1
0	0	1
1	4	5

and the two products are, with the element (i, j) of A as the row and the copy k as the column,

In CuTe: $A \otimes B$

	0	1	2	3	4	5
(0,0)	0	2	8	10	16	18
(0,1)	1	3	9	11	17	19
(1,0)	4	6	12	14	20	22
(1,1)	5	7	13	15	21	23

In our algebra: $\text{repeat}(A, B)$

	0	1	2	3	4	5
(0,0)	0	6	12	18	24	30
(0,1)	1	7	13	19	25	31
(1,0)	4	10	16	22	28	34
(1,1)	5	11	17	23	29	35

In CuTe the copy offsets are $A_{24}^* = (2, 3):(2, 8)$, that is 0, 2, 8, 10, 16, 18, and

$$A \otimes B = ((2, 2), (2, 3)) : ((4, 1), (2, 8))$$

is a bijection onto $[0, 24)$. In our algebra the copy offsets are the multiples of $\text{cosize}A = 6$, and $\text{repeat}(A, B) = (6, (2, 2)) : (6, (4, 1))$ is injective and not dense. CuTe’s placement is available in our algebra: $\text{divide}(\text{canonical}(R), S)$ is the pair (A, A_M^*) of the complement paragraph, and composing it on the right with $\text{repeat}(B, \text{canonical}(S_A))$, which sends (a, b) to $(a, B(b))$, gives $A \otimes B$ for every dense bijection B ; for $B = k:1$ that composition is the identity, and for this A

$$A \otimes B = \text{divide}(\text{canonical}((6, 4)), (2, 2)).$$

All of this is checked. With the operands exchanged and B a dense bijection, $B \otimes A = (B, S_A : (\text{size}B) D_A)$, which is $\text{interleave}(A, B)$ up to the order of the two coordinate groups. CuTe’s `blocked_product` and `raked_product` are both $A \otimes B$ with the two coordinate groups zipped per mode, block index first or copy index first: the same map on coordinates, regrouped. For $A = 4:1$ and $B = 2:1$, $A \otimes B = \text{repeat}(A, B) = (2, 4):(4, 1)$, `raked_product` is $(4, 2):(1, 4)$, the same map with its two coordinates exchanged, and $B \otimes A = \text{interleave}(A, B) = (2, 4):(1, 2)$.

Conditions. CuTe’s six conditions exist because its composition and complement must return a layout, so the operands are checked in advance for the result to be expressible as one. In our algebra neither operation has to: a composite is a map, and whether it is a layout is decided afterwards, for code generation only; and a complement is never formed on its own, it is the rest of a divide. So none of the six is checked. What is checked instead is that the right operand of a composition is a dense bijection onto the left operand’s coordinates, by the stride criterion of Lemma 1; that a map into logical or thread-value space is a dense bijection and a map into physical space is alias-free; and the precondition of each operation, such as $k_i \mid n_i$ for divide.

Example. Two hardware instructions and a kernel’s use of them. The thread-value maps are transcribed from the PTX ISA fragment tables [8]. For `mma.m16n8k16`, lane $4g + t_4$ holds four accumulator values (v_m, v_n) , at row $g + 8v_m$ and column $2t_4 + v_n$ of the 16×8 tile; with the tile row-major,

$$F_C = ((8, 4), (2, 2)) : ((8, 2), (64, 1)) : \text{thread-value} \rightarrow \text{logical}.$$

For `ldmatrix`, lane $4g + q$ holds elements $(g, 2q)$ and $(g, 2q + 1)$ of an 8×8 matrix,

$$F_{\text{ld}} = ((8, 4), 2) : ((8, 2), 1) : \text{thread-value} \rightarrow \text{logical}.$$

Printed as CuTe prints thread-value layouts, one cell per element of the tile, labelled with the lane and value that hold it and coloured by lane:

F_C : the 16×8 accumulator

	0	1	2	3	4	5	6	7
0	T0 V0	T0 V1	T1 V0	T1 V1	T2 V0	T2 V1	T3 V0	T3 V1
1	T4 V0	T4 V1	T5 V0	T5 V1	T6 V0	T6 V1	T7 V0	T7 V1
2	T8 V0	T8 V1	T9 V0	T9 V1	T10 V0	T10 V1	T11 V0	T11 V1
3	T12 V0	T12 V1	T13 V0	T13 V1	T14 V0	T14 V1	T15 V0	T15 V1
4	T16 V0	T16 V1	T17 V0	T17 V1	T18 V0	T18 V1	T19 V0	T19 V1
5	T20 V0	T20 V1	T21 V0	T21 V1	T22 V0	T22 V1	T23 V0	T23 V1
6	T24 V0	T24 V1	T25 V0	T25 V1	T26 V0	T26 V1	T27 V0	T27 V1
7	T28 V0	T28 V1	T29 V0	T29 V1	T30 V0	T30 V1	T31 V0	T31 V1
8	T0 V2	T0 V3	T1 V2	T1 V3	T2 V2	T2 V3	T3 V2	T3 V3
9	T4 V2	T4 V3	T5 V2	T5 V3	T6 V2	T6 V3	T7 V2	T7 V3
10	T8 V2	T8 V3	T9 V2	T9 V3	T10 V2	T10 V3	T11 V2	T11 V3
11	T12 V2	T12 V3	T13 V2	T13 V3	T14 V2	T14 V3	T15 V2	T15 V3
12	T16 V2	T16 V3	T17 V2	T17 V3	T18 V2	T18 V3	T19 V2	T19 V3
13	T20 V2	T20 V3	T21 V2	T21 V3	T22 V2	T22 V3	T23 V2	T23 V3
14	T24 V2	T24 V3	T25 V2	T25 V3	T26 V2	T26 V3	T27 V2	T27 V3
15	T28 V2	T28 V3	T29 V2	T29 V3	T30 V2	T30 V3	T31 V2	T31 V3

F_{ld} : one 8×8 matrix

	0	1	2	3	4	5	6	7
0	T0 V0	T0 V1	T1 V0	T1 V1	T2 V0	T2 V1	T3 V0	T3 V1
1	T4 V0	T4 V1	T5 V0	T5 V1	T6 V0	T6 V1	T7 V0	T7 V1
2	T8 V0	T8 V1	T9 V0	T9 V1	T10 V0	T10 V1	T11 V0	T11 V1
3	T12 V0	T12 V1	T13 V0	T13 V1	T14 V0	T14 V1	T15 V0	T15 V1
4	T16 V0	T16 V1	T17 V0	T17 V1	T18 V0	T18 V1	T19 V0	T19 V1
5	T20 V0	T20 V1	T21 V0	T21 V1	T22 V0	T22 V1	T23 V0	T23 V1
6	T24 V0	T24 V1	T25 V0	T25 V1	T26 V0	T26 V1	T27 V0	T27 V1
7	T28 V0	T28 V1	T29 V0	T29 V1	T30 V0	T30 V1	T31 V0	T31 V1

A kernel uses them through the operations. Two warps, $W = 2:1$, computing a 16×16 tile side by side have the thread-value map

$$\text{divide}(\text{canonical}((16, 16)), (16, 8)) \circ \text{interleave}(F_C, W),$$

with domain (warp, (lane, value)), and one warp loading the four 8×8 matrices of an 8×32 tile with `ldmatrix.x4` has

$$\text{divide}(\text{canonical}((8, 32)), (8, 8)) \circ \text{interleave}(F_{ld}, 4:1),$$

with domain (matrix, (lane, value)). Coloured by warp and by matrix respectively; in the first, cells are labelled by the thread $32w + \text{lane}$:

two warps on 16×16

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	T0 V0	T0 V1	T1 V0	T1 V1	T2 V0	T2 V1	T3 V0	T3 V1	T32 V0	T32 V1	T33 V0	T33 V1	T34 V0	T34 V1	T35 V0	T35 V1
1	T4 V0	T4 V1	T5 V0	T5 V1	T6 V0	T6 V1	T7 V0	T7 V1	T36 V0	T36 V1	T37 V0	T37 V1	T38 V0	T38 V1	T39 V0	T39 V1
2	T8 V0	T8 V1	T9 V0	T9 V1	T10 V0	T10 V1	T11 V0	T11 V1	T40 V0	T40 V1	T41 V0	T41 V1	T42 V0	T42 V1	T43 V0	T43 V1
3	T12 V0	T12 V1	T13 V0	T13 V1	T14 V0	T14 V1	T15 V0	T15 V1	T44 V0	T44 V1	T45 V0	T45 V1	T46 V0	T46 V1	T47 V0	T47 V1
4	T16 V0	T16 V1	T17 V0	T17 V1	T18 V0	T18 V1	T19 V0	T19 V1	T48 V0	T48 V1	T49 V0	T49 V1	T50 V0	T50 V1	T51 V0	T51 V1
5	T20 V0	T20 V1	T21 V0	T21 V1	T22 V0	T22 V1	T23 V0	T23 V1	T52 V0	T52 V1	T53 V0	T53 V1	T54 V0	T54 V1	T55 V0	T55 V1
6	T24 V0	T24 V1	T25 V0	T25 V1	T26 V0	T26 V1	T27 V0	T27 V1	T56 V0	T56 V1	T57 V0	T57 V1	T58 V0	T58 V1	T59 V0	T59 V1
7	T28 V0	T28 V1	T29 V0	T29 V1	T30 V0	T30 V1	T31 V0	T31 V1	T60 V0	T60 V1	T61 V0	T61 V1	T62 V0	T62 V1	T63 V0	T63 V1
8	T0 V2	T0 V3	T1 V2	T1 V3	T2 V2	T2 V3	T3 V2	T3 V3	T32 V2	T32 V3	T33 V2	T33 V3	T34 V2	T34 V3	T35 V2	T35 V3
9	T4 V2	T4 V3	T5 V2	T5 V3	T6 V2	T6 V3	T7 V2	T7 V3	T36 V2	T36 V3	T37 V2	T37 V3	T38 V2	T38 V3	T39 V2	T39 V3
10	T8 V2	T8 V3	T9 V2	T9 V3	T10 V2	T10 V3	T11 V2	T11 V3	T40 V2	T40 V3	T41 V2	T41 V3	T42 V2	T42 V3	T43 V2	T43 V3
11	T12 V2	T12 V3	T13 V2	T13 V3	T14 V2	T14 V3	T15 V2	T15 V3	T44 V2	T44 V3	T45 V2	T45 V3	T46 V2	T46 V3	T47 V2	T47 V3
12	T16 V2	T16 V3	T17 V2	T17 V3	T18 V2	T18 V3	T19 V2	T19 V3	T48 V2	T48 V3	T49 V2	T49 V3	T50 V2	T50 V3	T51 V2	T51 V3
13	T20 V2	T20 V3	T21 V2	T21 V3	T22 V2	T22 V3	T23 V2	T23 V3	T52 V2	T52 V3	T53 V2	T53 V3	T54 V2	T54 V3	T55 V2	T55 V3
14	T24 V2	T24 V3	T25 V2	T25 V3	T26 V2	T26 V3	T27 V2	T27 V3	T56 V2	T56 V3	T57 V2	T57 V3	T58 V2	T58 V3	T59 V2	T59 V3
15	T28 V2	T28 V3	T29 V2	T29 V3	T30 V2	T30 V3	T31 V2	T31 V3	T60 V2	T60 V3	T61 V2	T61 V3	T62 V2	T62 V3	T63 V2	T63 V3

ldmatrix.x4 on 8×32

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
0	T0 V0	T0 V1	T1 V0	T1 V1	T2 V0	T2 V1	T3 V0	T3 V1	T0 V0	T0 V1	T1 V0	T1 V1	T2 V0	T2 V1	T3 V0	T3 V1	T0 V0	T0 V1	T1 V0	T1 V1	T2 V0	T2 V1	T3 V0	T3 V1	T0 V0	T0 V1	T1 V0	T1 V1	T2 V0	T2 V1	T3 V0	T3 V1
1	T4 V0	T4 V1	T5 V0	T5 V1	T6 V0	T6 V1	T7 V0	T7 V1	T4 V0	T4 V1	T5 V0	T5 V1	T6 V0	T6 V1	T7 V0	T7 V1	T4 V0	T4 V1	T5 V0	T5 V1	T6 V0	T6 V1	T7 V0	T7 V1	T4 V0	T4 V1	T5 V0	T5 V1	T6 V0	T6 V1	T7 V0	T7 V1
2	T8 V0	T8 V1	T9 V0	T9 V1	T10 V0	T10 V1	T11 V0	T11 V1	T8 V0	T8 V1	T9 V0	T9 V1	T10 V0	T10 V1	T11 V0	T11 V1	T8 V0	T8 V1	T9 V0	T9 V1	T10 V0	T10 V1	T11 V0	T11 V1	T8 V0	T8 V1	T9 V0	T9 V1	T10 V0	T10 V1	T11 V0	T11 V1
3	T12 V0	T12 V1	T13 V0	T13 V1	T14 V0	T14 V1	T15 V0	T15 V1	T12 V0	T12 V1	T13 V0	T13 V1	T14 V0	T14 V1	T15 V0	T15 V1	T12 V0	T12 V1	T13 V0	T13 V1	T14 V0	T14 V1	T15 V0	T15 V1	T12 V0	T12 V1	T13 V0	T13 V1	T14 V0	T14 V1	T15 V0	T15 V1
4	T16 V0	T16 V1	T17 V0	T17 V1	T18 V0	T18 V1	T19 V0	T19 V1	T16 V0	T16 V1	T17 V0	T17 V1	T18 V0	T18 V1	T19 V0	T19 V1	T16 V0	T16 V1	T17 V0	T17 V1	T18 V0	T18 V1	T19 V0	T19 V1	T16 V0	T16 V1	T17 V0	T17 V1	T18 V0	T18 V1	T19 V0	T19 V1
5	T20 V0	T20 V1	T21 V0	T21 V1	T22 V0	T22 V1	T23 V0	T23 V1	T20 V0	T20 V1	T21 V0	T21 V1	T22 V0	T22 V1	T23 V0	T23 V1	T20 V0	T20 V1	T21 V0	T21 V1	T22 V0	T22 V1	T23 V0	T23 V1	T20 V0	T20 V1	T21 V0	T21 V1	T22 V0	T22 V1	T23 V0	T23 V1
6	T24 V0	T24 V1	T25 V0	T25 V1	T26 V0	T26 V1	T27 V0	T27 V1	T24 V0	T24 V1	T25 V0	T25 V1	T26 V0	T26 V1	T27 V0	T27 V1	T24 V0	T24 V1	T25 V0	T25 V1	T26 V0	T26 V1	T27 V0	T27 V1	T24 V0	T24 V1	T25 V0	T25 V1	T26 V0	T26 V1	T27 V0	T27 V1
7	T28 V0	T28 V1	T29 V0	T29 V1	T30 V0	T30 V1	T31 V0	T31 V1	T28 V0	T28 V1	T29 V0	T29 V1	T30 V0	T30 V1	T31 V0	T31 V1	T28 V0	T28 V1	T29 V0	T29 V1	T30 V0	T30 V1	T31 V0	T31 V1	T28 V0	T28 V1	T29 V0	T29 V1	T30 V0	T30 V1	T31 V0	T31 V1

Both are dense bijections, as Definition 7 requires of a map into logical. Composed onto storage they become address maps, and Section 4 emits their code. With the 16×16 tile stored row-major with leading dimension 24, that is with $(16, 16):(24, 1)$: logical $\rightarrow$ physical in place of canonical $((16, 16))$, the two-warp map is decided to be a layout and the emitted address is

$$8*c0 + 24*c1_0_0 + 2*c1_0_1 + 192*c1_1_0 + c1_1_1$$

where $c0$ is the warp and $c1_0_0$, $c1_0_1$, $c1_1_0$, $c1_1_1$ are g , t_4 , v_m , v_n : the address $24(g + 8v_m) + 8w + 2t_4 + v_n$. The slice for warp 1 (Section 5) is

$$24*c1_0_0 + 2*c1_0_1 + 192*c1_1_0 + c1_1_1 + 8$$

With the 8×32 tile stored with leading dimension 40, the ldmatrix map emits

$$8*c0 + 40*c1_0_0 + 2*c1_0_1 + c1_1$$

with $c0$ the matrix and $c1_0_0$, $c1_0_1$, $c1_1$ the lane's g and q and the value v . The tests build the same construction for a 32×16 tile and four warps, and the ldmatrix reader of a GEMM mainloop (Section 7).

Usage. The table lists the idioms CUTLASS kernels are written with and what each is in our algebra. $L|_{x=v}$ fixes the entries named x to v (Definition 9). F : thread-value $\rightarrow$ logical is a valid map, such as F_C above; W is a dense bijection arranging the warps, as above; and T : logical $\rightarrow$ thread-value is a dense bijection from the positions of a tile to the threads.

CuTe	Our algebra
<code>make_layout(S, D)</code>	$S:D$, modes reversed
<code>composition(A, B)</code>	$A \circ B$
<code>complement(B, M)</code>	rest of $\text{divide}(\text{canonical}(R), S)$, above
<code>logical_divide(A, S),</code> <code>zipped_divide</code>	$\text{divide}(A, S)$
<code>local_tile(A, S, b)</code>	$\text{divide}(A, S) _{\text{rest}=b}$
<code>local_partition(A, T, t)</code>	$\text{divide}(A, S_T) _{\text{tile}=\text{inverse}(T)(t)}$
<code>partition_A/B/C, tensor(_, k)</code>	$(A \circ F) _{\text{thread}=t}; A _k$
<code>logical_product(A, B),</code> <code>blocked_product, tile_to_shape</code>	$\text{repeat}(A, B)$, A a dense bijection
<code>raked_product(A, B)</code>	$\text{repeat}(A, B)$ with the two coordinate groups exchanged in each mode
<code>logical_product(B, A),</code> copies first	$\text{interleave}(A, B)$, B a dense bijection, up to the order of the groups
stride-0 modes	$\text{broadcast}(L, S)$; read-only
<code>Swizzle<B,M,S> \circ L</code>	$\sigma \circ L$ with $b = B$, $dst = M$, $src = M+S$ (Definition 4)
k -stage pipeline buffer	$\text{repeat}(L, k:1)$
<code>TiledCopy: partition_S/D, vector</code> width v	$\text{divide}(A, (1, v)) \circ \text{interleave}(v:1, 32:1)$ over the lanes
<code>TiledMMA: F</code> over a CTA tile with warp grid W	$\text{divide}(A, S_F) \circ \text{interleave}(F, W)$
<code>right_inverse(L); retile_S/D</code>	$\text{inverse}(L); \text{inverse}(F_{\text{mma}}) \circ F_{\text{copy}}$
<code>coalesce</code>	not needed: the kernel declares the shape it iterates
<code>group_modes, select, flatten</code>	regrouping the shape; same function
<code>recast, upcast by v</code>	$\text{divide}(A, (1, \dots, 1, v))$; the change of element size is outside the algebra
TMA box S ; multicast	$\text{divide}(A, S)$; broadcast across CTAs

Every row is one of composition, divide, restrict, repeat, interleave, broadcast, a swizzle, inverse, or a regrouping of a shape.

7 Evaluation

The implementation¹ is 956 lines of OCaml, over 403 lines of interface; the tests are 2,432. Every figure quoted below is generated from a run of the test suite (`paper/eval.sh`), not transcribed: the random generators are seeded, so the numbers are reproducible on the tree that produced them.

The decision procedure is correct. The criterion of Lemma 1 agrees with enumeration on 400 random layouts. The scan of Theorem 1 is compared with an oracle that tries every chain of divisors of n and checks the strides it forces at every point. Over every function

¹<https://github.com/annp0/simplified-layout>

$g : [0, n) \rightarrow [0, k)$ with $g(0) = 0$, for $n \in \{2, 3, 4, 6, 8, 9, 12\}$ and $k = 5$ if $n \leq 4$, 3 otherwise, that is 186,293 functions, the two accept the same 58, with the same chain.

CuTe’s conditions are not checked, and the failure is silent. The sweep quoted in Section 1 takes every rank-2 layout $(N_0, N_1):(d_0, d_1)$ with $N_0, N_1 \in \{2, 3, 4\}$, $d_0 \in \{1, 2\}$ and $d_1 \leq 16$, against cotargets $M \in \{12, 24, 48\}$, keeping those for which the complement formula is defined at all — CuTe rejects $d_1 < N_0 d_0$ with a `static_assert`, the one condition it does enforce. Every entry is a static integer, so nothing here is hidden from the compiler. Each complement is checked against the postconditions CuTe’s own unit test asserts: that $\text{cosize}(A, A^*)$ reach M , that A^* be ordered, and that A^* and A be disjoint. Cases are classified by *both* of the complement’s divisibility conditions, $N_0 d_0 \mid d_1$ and $N_1 d_1 \mid M$. Of the 675 cases, the 78 satisfying both all pass; of the 597 that violate one or the other, 40 fail. The conditions are therefore sufficient and not necessary, which is exactly why their absence bites: most violations happen to be harmless, nothing distinguishes those from the rest, and the ones that are not harmless fail the codomain bound — never orderedness, never disjointness — without a diagnostic. The mechanism is plain in the formula: its shape entries are quotients, so when a divisibility fails the integer division truncates and the complement simply covers less than it should. For $(2, 2):(1, 3)$ at $M = 24$ the entry $d_1/N_0 d_0 = 3/2$ becomes 1, and cosize comes out 23. Measured against CUTLASS 4.8; the harness is `study/run.sh` and needs no GPU.

The decision procedure is machine-checked. Lemma 1, Lemma 2 and Theorem 1 are proved in Coq, 1,967 lines, each as an *if and only if*; no proof is admitted and `Print Assumptions` reports every one closed under the global context. Composing them gives the statement about the procedure as a whole: the checks of this section — separability in one pass over the box, then the scan of Theorem 1 on each axis — succeed exactly when the map is the index function of a layout over a refinement of its domain, plus a constant. One direction is what stops a wrong address expression being emitted; the other is what makes “decided rather than searched” true, since it says no simplification is missed. Two things are proved differently from the presentation above. The scan is phrased with its residual recomputed from the weights committed so far rather than by mutating a copy of the first difference: that is the invariant the proof of Theorem 1 establishes anyway, so it is the same number at every step. And the forward direction of Lemma 1 avoids the counting step of its proof — that modes $1, \dots, j-1$ produce at most N_{j-1} values — in favour of the observation that disjoint blocks of n consecutive positions covering $[0, M)$ must be the aligned ones, where the covering comes from surjectivity and the disjointness from injectivity, so no cardinality argument is needed. The proved scan and the implementation are not merely alike: the scan is extracted from Coq to OCaml and run beside `fit_axis` on every one of the 186,293 maps of the enumeration above, where they agree on the verdict and, on the 58 they accept, digit for digit. The extraction is faithful — no datatype is remapped onto a machine integer — so nothing is assumed of it that the proofs do not establish. Not mechanized: the minimality of the coarsest chain, which the tests check against the brute-force oracle; the regrouping of a nested shape into the flat shape of its leaves; and the reordering of modes by stride, the proofs being stated for modes already sorted.

It matters in practice. The tests build 207 layouts — CuTe’s documented examples, CUTLASS’s atoms, the usage table of Section 6, and the laws of Section 3 — and every one of them is checked against its emitted expression at every coordinate. Of those, 172 are affine in their

coordinates, 11 are layouts of a refinement with one coordinate split, and 24 are not layouts and are emitted as composites. On 600 random composites of a dense bijection with a padded storage map, 358 are affine and 531 are layouts of some refinement, and their emitted forms contain 646 divisions and moduli where the composites contain 976. Chaining two logical stages onto storage rather than one, 242 of 300 are still layouts of some refinement, 115 of them affine: the property survives composition. With a random swizzle on the storage map none of 300 is a layout, though all 300 of the same composites are without it — the swizzle alone is what costs the strided form.

The algebra is enough for a kernel. The SM90 GMMA 64×128 accumulator, built from the SM80 `mma.m16n8k16 C` fragment as

$$\text{divide}(\text{canonical}((64, 128)), (16, 8)) \circ \text{interleave}(F_C, \text{canonical}((4, 16))),$$

equals NVIDIA’s printed layout at all 8,192 coordinates. A GEMM mainloop built from the operations alone, an 8×32 tile of 16-bit elements copied in 16-byte vectors by one warp into two shared-memory stages of the SW64 atom and read back by `ldmatrix`, with the writer and reader of Section 6 and the stages by repeat, is consistent: for every element, the address the copy writes is the address the `mma` operand is read from.

8 Related work

CuTe [1, 2] is the system under study; its operations and their conditions are stated in Section 1. **Shah** [3] formalizes the algebra and gives the conditions under which the composite of two layouts is again a layout. In our algebra that is decided after the fact by Theorem 1, and a composite that is not a layout is still a legal map. **Carlisle et al.** [4] model a class of CuTe layouts as morphisms of two categories, define composition, logical product and logical division on the morphisms, and characterize the layouts so obtained. **Linear layouts** [5], in Triton, represent a layout as a matrix over $\mathbb{F}_2$ acting on the bits of the hardware index, which gives generic layout-to-layout conversion; a swizzle of Definition 4 is such a matrix. **Bhaskaracharya et al.** [6] represent both CuTe layouts and linear layouts as integer set relations in ISL and implement composition, inversion and complement with its operations; the question of Section 4 can be posed there, and Theorem 1 answers it without a solver. **Axe** [7] maps logical coordinates to a multi-axis physical space through named axes and unifies tiling, sharding, replication and offsets from device meshes to threads; its replication is explicit, as broadcast is in our algebra, and its Appendix A.2.1 uses the floor expansion and first differences of Section 4 to prove a normalized representation unique. **Halide** [9], **TVM** [10] and **MLIR** [11] carry general integer expression simplifiers; Section 4 decides the layout question directly instead.

9 The DSL this serves

The algebra is built for a GPU DSL in which each warp is a SIMD program and warps communicate through pipes, in the sense of Unix pipes, lowered to mbarriers at compile time. The layouts of this paper describe how the lanes of a warp access a tile.

References

- [1] NVIDIA. *CuTe: Layouts and Layout Algebra*. CUTLASS documentation, docs.nvidia.com/cutlass, accessed 2026.
- [2] C. Cecka. *CuTe Layout Representation and Algebra*. NVIDIA Research; arXiv:2603.02298, 2026.
- [3] J. Shah. *A Note on the Algebra of CuTe Layouts*. Colfax Research technical note, 2024.
- [4] J. Carlisle, J. Shah, R. Stern, P. VanKoughnett. *Categorical Foundations for CuTe Layouts*. arXiv:2601.05972, 2026.
- [5] K. Zhou, M. Lezcano, A. Goucher, A. Rakhmati, J. Niu, J. Lebar, P. Szczerbuk, P. Bell, P. Tillet, T. Raoux, Z. Moudallal. *Linear Layouts: Robust Code Generation of Efficient Tensor Computation Using $\mathbb{F}_2$* . ASPLOS 2026; arXiv:2505.23819.
- [6] S. G. Bhaskaracharya, A. Acharya, B. Hagedorn, V. Grover. *Modeling Layout Abstractions Using Integer Set Relations*. arXiv:2511.10374, 2025.
- [7] B. Hou, H. Jin, G. Wang, J. Chen, Y. Cai, L. Yang, Z. Ye, Y. Ding, R. Lai, T. Chen. *Axe: A Simple Unified Layout Abstraction for Machine Learning Compilers*. arXiv:2601.19092, 2026.
- [8] NVIDIA. *Parallel Thread Execution ISA*. Matrix fragment layout tables, accessed 2026.
- [9] J. Ragan-Kelley, C. Barnes, A. Adams, S. Paris, F. Durand, S. Amarasinghe. *Halide: A Language and Compiler for Optimizing Parallelism, Locality, and Recomputation in Image Processing Pipelines*. PLDI 2013.
- [10] T. Chen et al. *TVM: An Automated End-to-End Optimizing Compiler for Deep Learning*. OSDI 2018.
- [11] C. Lattner et al. *MLIR: Scaling Compiler Infrastructure for Domain Specific Computation*. CGO 2021.